

10. Parametric Inference

Parametric models are of the form

$$\mathfrak{F} = \left\{ f(x; \theta) : \theta \in \Theta \right\}$$

where $\Theta \subset \mathbb{R}^k$ is the parameter space and $\theta = (\theta_1, \dots, \theta_k)$ is the parameter. The problem of inference then reduces to the problem of estimating parameter θ .

10.1 Parameter of interest

Often we are only interested in some function $T(\theta)$. For example, if $X \sim N(\mu, \sigma^2)$ then the parameter is $\theta = (\mu, \sigma)$. If our goal is to estimate μ then $\mu = T(\theta)$ is called the **parameter of interest** and σ is called a **nuisance parameter**.

10.2 The Method of Moments

Suppose that the parameter $\theta = (\theta_1, \dots, \theta_n)$ has k components. For $1 \leq j \leq k$ define the j -th **moment**

$$\alpha_j \equiv \alpha_j(\theta) = \mathbb{E}_\theta(X^j) = \int x^j dF_\theta(x)$$

and the j -th **sample moment**

$$\hat{\alpha}_j = \frac{1}{n} \sum_{i=1}^n X_i^j$$

The **method of moments estimator** $\hat{\theta}_n$ is defined to be the value of θ such that

$$\alpha_1(\hat{\theta}_n) = \hat{\alpha}_1 \tag{1}$$

$$\alpha_2(\hat{\theta}_n) = \hat{\alpha}_2 \tag{2}$$

$$\vdots \tag{3}$$

$$\alpha_k(\hat{\theta}_n) = \hat{\alpha}_k \tag{4}$$

This defines a system of k equations with k unknowns.

Theorem 10.6. Let $\hat{\theta}_n$ denote the method of moments estimator. Under the conditions given in the appendix, the following statements hold:

(1) The estimate $\hat{\theta}_n$ exists with probability tending to 1.

(2) The estimate is consistent: $\hat{\theta}_n \xrightarrow{P} \theta$.

(3) The estimate is asymptotically Normal:

$$\sqrt{n}(\hat{\theta}_n - \theta) \rightsquigarrow N(0, \Sigma)$$

where

$$\Sigma = g \mathbb{E}_\theta(YY^T)g^T$$

$$Y = (X, X^2, \dots, X^k)^T, \quad g = (g_1, \dots, g_k) \quad \text{and} \quad g_j = \partial \alpha_j^{-1}(\theta) / \partial \theta$$

The last statement in Theorem 10.6 can be used to find standard errors and confidence intervals. However, there is an easier way: the bootstrap.

10.3 Maximum Likelihood

Let X_1, \dots, X_n be iid with PDF $f(x; \theta)$.

The **likelihood function** is defined by

$$\mathcal{L}_n(\theta) = \prod_{i=1}^n f(X_i; \theta)$$

The **log-likelihood function** is defined by $\ell_n(\theta) = \log \mathcal{L}_n(\theta)$.

The likelihood function is just the joint density of the data, except we treat it as a function of parameter θ . Thus $\mathcal{L}_n : \Theta \rightarrow [0, \infty)$. The likelihood function is not a density function; in general it is not true that \mathcal{L}_n integrates to 1.

The **maximum likelihood estimator** MLE, denoted by $\hat{\theta}_n$, is the value of θ that maximizes $\mathcal{L}_n(\theta)$.

The maximum of $\ell_n(\theta)$ occurs at the same place as the maximum of $\mathcal{L}_n(\theta)$, so maximizing either leads to the same answer. Often it's easier to maximize the log-likelihood.

10.4 Properties of Maximum Likelihood Estimators

Under certain conditions on the model, the MLE $\hat{\theta}_n$ possesses many properties that make it an appealing choice of estimator.

The main properties of the MLE are:

- It is **consistent**: $\hat{\theta}_n \xrightarrow{P} \theta_*$, where θ_* denotes the true value of parameter θ .
- It is **equivariant**: if $\hat{\theta}_n$ is the MLE of θ then $g(\hat{\theta}_n)$ is the MLE of $g(\theta)$.
- It is **asymptotically Normal**: $\sqrt{n}(\hat{\theta} - \theta_*)/\hat{se} \rightsquigarrow N(0, 1)$ where \hat{se} can be computed analytically.
- It is **asymptotically optimal** or **efficient**: roughly, this means that among all well behaved estimators, the MLE has the smallest variance, at least for large samples.
- The MLE is approximately the Bayes estimator.

10.5 Consistency of Maximum Likelihood Estimator

If f and g are PDFs, define the **Kullback-Leibler distance** between f and g to be:

$$D(f, g) = \int f(x) \log \left(\frac{f(x)}{g(x)} \right) dx$$

It can be shown that $D(f, g) \geq 0$ and $D(f, f) = 0$. For any $\theta, \psi \in \Theta$ write $D(\theta, \psi)$ to mean $D(f(x; \theta), f(x; \psi))$. We will assume that $\theta \neq \psi$ implies $D(\theta, \psi) > 0$.

Let θ_* denote the true value of θ . Maximizing $\ell_n(\theta)$ is equivalent to maximizing

$$M_n(\theta) = \frac{1}{n} \sum_i \log \frac{f(X_i; \theta)}{f(X_i; \theta_*)}$$

By the law of large numbers, $M_n(\theta)$ converges to:

$$\mathbb{E}_{\theta_*} \left(\log \frac{f(X_i; \theta)}{f(X_i; \theta_*)} \right) = \int \log \left(\frac{f(x; \theta)}{f(x; \theta_*)} \right) f(x; \theta_*) dx \quad (5)$$

$$= - \int \log \left(\frac{f(x; \theta_*)}{f(x; \theta)} \right) f(x; \theta_*) dx \quad (6)$$

$$= -D(\theta_*, \theta) \quad (7)$$

Hence $M_n(\theta) \approx -D(\theta_*, \theta)$ which is maximized at θ_* , since the KL distance is 0 when $\theta_* = \theta$ and positive otherwise. Hence, we expect that the maximizer will tend to θ_* .

To prove this formally, we need more than $M_n(\theta) \xrightarrow{P} -D(\theta_*, \theta)$. We need this convergence to be uniform over θ . We also have to make sure that the KL distance is well-behaved. Here are the formal details.

Theorem 10.13. Let θ_* denote the true value of θ . Define

$$M_n(\theta) = \frac{1}{n} \sum_i \log \frac{f(X_i; \theta)}{f(X_i; \theta_*)}$$

and $M(\theta) = -D(\theta_*, \theta)$. Suppose that

$$\sup_{\theta \in \Theta} |M_n(\theta) - M(\theta)| \xrightarrow{P} 0$$

and that, for every $\epsilon > 0$,

$$\sup_{\theta: |\theta - \theta_*| \geq \epsilon} M(\theta) < M(\theta_*)$$

Let $\hat{\theta}_n$ denote the mle. Then $\hat{\theta}_n \xrightarrow{P} \theta_*$.

10.6 Equivalence of the MLE

Theorem 10.14. Let $\tau = g(\theta)$ be a one-to-one function of θ . Let $\hat{\theta}_n$ be the MLE of θ . Then $\hat{\tau}_n = g(\hat{\theta}_n)$ is the MLE of τ .

Proof. Let $h = g^{-1}$ denote the inverse of g . Then $\hat{\theta}_n = h(\hat{\tau}_n)$. For any τ , $L(\tau) = \prod_i f(x_i; h(\tau)) = \prod_i f(x_i; \theta) = \mathcal{L}(\theta)$ where $\theta = h(\tau)$. Hence, for any τ , $\mathcal{L}_n(\tau) = \mathcal{L}(\theta) \leq \mathcal{L}(\hat{\theta}) = \mathcal{L}_n(\hat{\tau})$.

10.7 Asymptotic Normality

The **score function** is defined to be

$$s(X; \theta) = \frac{\partial \log f(X; \theta)}{\partial \theta}$$

The **Fisher information** is defined to be

$$I_n(\theta) = \mathbb{V}_\theta \left(\sum_{i=1}^n s(X_i; \theta) \right) \quad (8)$$

$$= \sum_{i=1}^n \mathbb{V}_\theta(s(X_i; \theta)) \quad (9)$$

For $n = 1$ we sometimes write $I(\theta)$ instead of $I_1(\theta)$.

It can be shown that $\mathbb{E}_\theta(s(X; \theta)) = 0$. It then follows that $\mathbb{V}_\theta(s(X; \theta)) = \mathbb{E}_\theta((s(X; \theta))^2)$. A further simplification of $I_n(\theta)$ is given in the next result.

Theorem 10.17.

$$I_n(\theta) = nI(\theta)$$

$$I(\theta) = -\mathbb{E}_\theta \left(\frac{\partial^2 \log f(X; \theta)}{\partial^2 \theta^2} \right) \quad (10)$$

$$= -\int \left(\frac{\partial^2 \log f(x; \theta)}{\partial^2 \theta^2} \right) f(x; \theta) dx \quad (11)$$

Theorem 10.18 (Asymptotic Normality of the MLE). Under appropriate regularity conditions, the following hold:

(1) Let $se = \sqrt{1/I_n(\theta)}$. Then,

$$\frac{\hat{\theta}_n - \theta}{se} \rightsquigarrow N(0, 1)$$

(2) Let $\hat{se} = \sqrt{1/I_n(\hat{\theta}_n)}$. Then,

$$\frac{\hat{\theta}_n - \theta}{\hat{se}} \rightsquigarrow N(0, 1)$$

The first statement says that $\hat{\theta}_n \approx N(\theta, se)$. The second statement says that this is still true if we replace the standard error se by its estimated standard error \hat{se} .

Informally this says that the distribution of the MLE can be approximated with $N(\theta, \hat{se})$. From this fact we can construct an asymptotic confidence interval.

Theorem 10.19. Let

$$C_n = \left(\hat{\theta}_n - z_{\alpha/2} \hat{\text{se}}, \hat{\theta}_n + z_{\alpha/2} \hat{\text{se}} \right)$$

Then, $\mathbb{P}_\theta(\theta \in C_n) \rightarrow 1 - \alpha$ as $n \rightarrow \infty$.

Proof Let Z denote a standard random variable. Then,

$$\mathbb{P}_\theta(\theta \in C_n) = \mathbb{P}_\theta(\hat{\theta}_n - z_{\alpha/2} \hat{\text{se}} \leq \theta \leq \hat{\theta}_n + z_{\alpha/2} \hat{\text{se}}) \quad (12)$$

$$= \mathbb{P}_\theta(-z_{\alpha/2} \leq \frac{\hat{\theta}_n - \theta}{\hat{\text{se}}} \leq z_{\alpha/2}) \quad (13)$$

$$\rightarrow \mathbb{P}(-z_{\alpha/2} \leq Z \leq z_{\alpha/2}) = 1 - \alpha \quad (14)$$

10.8 Optimality

Suppose that $X_1, \dots, X_n \sim N(0, \sigma^2)$. The MLE is $\hat{\theta}_n = \bar{X}_n$. Another reasonable estimator is the sample median $\bar{\theta}_n$. The MLE satisfies

$$\sqrt{n}(\hat{\theta}_n - \theta) \rightsquigarrow N(0, \sigma^2)$$

It can be proved that the median satisfies

$$\sqrt{n}(\bar{\theta}_n - \theta) \rightsquigarrow N\left(0, \sigma^2 \frac{\pi}{2}\right)$$

This means that the median converges to the right value but has a larger variance than the MLE.

More generally, consider two estimators T_n and U_n and suppose that

$$\sqrt{n}(T_n - \theta) \rightsquigarrow N(0, t^2) \quad \text{and} \quad \sqrt{n}(U_n - \theta) \rightsquigarrow N(0, u^2)$$

We define the **asymptotic relative efficiency** of U to T by $ARE(U, T) = t^2/u^2$. In the Normal example, $ARE(\bar{\theta}_n, \hat{\theta}_n) = 2/\pi = 0.63$.

Theorem 10.23. If $\hat{\theta}_n$ is the MLE and $\bar{\theta}_n$ is any other estimator then

$$ARE(\bar{\theta}_n, \hat{\theta}_n) \leq 1$$

Thus, MLE has the smallest (asymptotic) variance and we say that MLE is **efficient** or **asymptotically optimal**.

The result is predicated over the model being correct -- otherwise the MLE may no longer be optimal.

10.9 The Delta Method

Let $\tau = g(\theta)$ where g is a smooth function. The maximum likelihood estimator of τ is $\hat{\tau} = g(\hat{\theta})$.

Theorem 10.24 (The Delta Method). If $\tau = g(\theta)$ where g is differentiable and $g'(\theta) \neq 0$ then

$$\frac{\sqrt{n}(\hat{\tau}_n - \tau)}{\hat{\text{se}}(\hat{\tau})} \rightsquigarrow N(0, 1)$$

where $\hat{\tau}_n = g(\hat{\theta})$ and

$$\hat{\text{se}}(\hat{\tau}_n) = |g'(\hat{\theta})| \hat{\text{se}}(\hat{\theta}_n)$$

Hence, if

$$C_n = (\hat{\tau}_n - z_{\alpha/2} \hat{\text{se}}(\hat{\tau}_n), \hat{\tau}_n + z_{\alpha/2} \hat{\text{se}}(\hat{\tau}_n))$$

then $\mathbb{P}_{\theta}(\tau \in C_n) \rightarrow 1 - \alpha$ as $n \rightarrow \infty$.

10.10 Multiparameter Models

We can extend these ideas to models with several parameters.

Let $\theta = (\theta_1, \dots, \theta_n)$ and let $\hat{\theta} = (\hat{\theta}_1, \dots, \hat{\theta}_n)$ be the MLE. Let $\ell_n = \sum_{i=1}^n \log f(X_i; \theta)$,

$$H_{jj} = \frac{\partial^2 \ell_n}{\partial \theta_j^2} \quad \text{and} \quad H_{jk} = \frac{\partial^2 \ell_n}{\partial \theta_j \partial \theta_k}$$

Define the **Fisher Information Matrix** by

$$I_n(\theta) = - \begin{bmatrix} \mathbb{E}_{\theta}(H_{11}) & \mathbb{E}_{\theta}(H_{12}) & \cdots & \mathbb{E}_{\theta}(H_{1k}) \\ \mathbb{E}_{\theta}(H_{21}) & \mathbb{E}_{\theta}(H_{22}) & \cdots & \mathbb{E}_{\theta}(H_{2k}) \\ \vdots & \vdots & \ddots & \vdots \\ \mathbb{E}_{\theta}(H_{k1}) & \mathbb{E}_{\theta}(H_{k2}) & \cdots & \mathbb{E}_{\theta}(H_{kk}) \end{bmatrix}$$

Let $J_n(\theta) = I_n^{-1}(\theta)$ be the inverse of I_n .

Theorem 10.27. Under appropriate regularity conditions,

$$\sqrt{n}(\hat{\theta} - \theta) \approx N(0, J_n(\theta))$$

Also, if $\hat{\theta}_j$ is the j -th component of $\hat{\theta}$, then

$$\frac{\sqrt{n}(\hat{\theta}_j - \theta_j)}{\hat{\text{se}}_j} \approx N(0, 1)$$

where $\hat{\text{se}}_j^2$ is the j -th diagonal element of J_n . The approximate covariance of $\hat{\theta}_j$ and $\hat{\theta}_k$ is $\text{Cov}(\hat{\theta}_j, \hat{\theta}_k) \approx J_n(j, k)$.

There is also a multiparameter delta method. Let $\tau = g(\theta_1, \dots, \theta_k)$ be a function and let

$$\nabla g = \begin{pmatrix} \frac{\partial g}{\partial \theta_1} \\ \vdots \\ \frac{\partial g}{\partial \theta_k} \end{pmatrix}$$

be the gradient of g .

Theorem 10.28 (Multiparameter delta method). Suppose that ∇g evaluated at $\hat{\theta}$ is not 0. Let $\hat{\tau} = g(\hat{\theta})$. Then

$$\frac{\sqrt{n}(\hat{\tau} - \tau)}{\hat{\text{se}}(\hat{\tau})} \rightsquigarrow N(0, 1)$$

where

$$\hat{\text{se}}(\hat{\tau}) = \sqrt{(\hat{\nabla} g)^T \hat{J}_n (\hat{\nabla} g)},$$

$\hat{J}_n = J_n(\hat{\theta}_n)$ and $\hat{\nabla} g$ is ∇g evaluated at $\theta = \hat{\theta}$.

10.11 The Parametric Bootstrap

For parametric models, standard errors and confidence intervals may also be estimated using the bootstrap. There is only one change. In nonparametric bootstrap, we sampled X_1^*, \dots, X_n^* from the empirical distribution \hat{F}_n . In the parametric bootstrap we sample instead from

$f(x; \hat{\theta}_n)$. Here, $\hat{\theta}_n$ could be the MLE or the method of moments estimator.

10.12 Technical Appendix

10.12.1 Proofs

Proof of Theorem 10.13. Since $\hat{\theta}_n$ maximizes $M_n(\theta)$, we have $M_n(\hat{\theta}) \geq M_n(\theta_*)$. Hence,

$$M(\theta_*) - M(\hat{\theta}_n) = M_n(\theta_*) - M(\hat{\theta}_n) + M(\hat{\theta}_n) - M_n(\theta_*) \quad (15)$$

$$\leq M_n(\hat{\theta}) - M(\hat{\theta}_n) + M(\theta_*) - M_n(\theta_*) \quad (16)$$

$$\leq \sup_{\theta} |M_n(\theta) - M(\theta)| + M(\theta_*) - M_n(\theta_*) \quad (17)$$

$$\xrightarrow{P} 0 \quad (18)$$

It follows that, for any $\delta > 0$,

$$\mathbb{P}(M(\hat{\theta}_n) < M(\theta_*) - \delta) \rightarrow 0$$

Pick any $\epsilon > 0$. There exists $\delta > 0$ such that $|\theta - \theta_*| \geq \epsilon$ implies that $M(\theta) < M(\theta_*) - \delta$. Hence,

$$\mathbb{P}(|\hat{\theta}_n - \theta_*| > \epsilon) \leq \mathbb{P}(M(\hat{\theta}_n) < M(\theta_*) - \delta) \rightarrow 0$$

Lemma 10.31. The score function satisfies

$$\mathbb{E}[s(X; \theta)] = 0$$

Proof. Note that $1 = \int f(x; \theta) dx$. Differentiate both sides of this equation to get

$$0 = \frac{\partial}{\partial \theta} \int f(x; \theta) dx = \int \frac{\partial}{\partial \theta} f(x; \theta) dx \quad (19)$$

$$= \int \frac{\frac{\partial f(x; \theta)}{\partial \theta}}{f(x; \theta)} f(x; \theta) dx = \int \frac{\partial \log f(x; \theta)}{\partial \theta} f(x; \theta) dx \quad (20)$$

$$= \int s(x; \theta) f(x; \theta) dx = \mathbb{E}[s(X; \theta)] \quad (21)$$

Proof of Theorem 10.18. Let $\ell(\theta) = \log \mathcal{L}(\theta)$. Then

$$0 = \ell'(\hat{\theta}) \approx \ell'(\theta) + (\hat{\theta} - \theta)\ell''(\theta)$$

Rearrange the above equation to get $\hat{\theta} - \theta = -\ell'(\theta)/\ell''(\theta)$, or

$$\sqrt{n}(\hat{\theta} - \theta) = \frac{\frac{1}{\sqrt{n}}\ell'(\theta)}{-\frac{1}{n}\ell''(\theta)} = \frac{\text{TOP}}{\text{BOTTOM}}$$

Let $Y_i = \partial \log f(X_i, \theta) / \partial \theta$. From the previous lemma $\mathbb{E}(Y_i) = 0$ and also $\mathbb{V}(Y_i) = I(\theta)$. Hence,

$$\text{TOP} = n^{-1/2} \sum_i Y_i = \sqrt{n} \bar{Y} = \sqrt{n}(\bar{Y} - 0) \rightsquigarrow W \sim N(0, I)$$

by the central limit theorem. Let $A_i = -\partial^2 \log f(X_i; \theta) / \partial \theta^2$. Then $\mathbb{E}(A_i) = I(\theta)$ and

$$\text{BOTTOM} = \bar{A} \xrightarrow{P} I(\theta)$$

by the law of large numbers. Apply Theorem 6.5 part (e) to conclude that

$$\sqrt{n}(\hat{\theta} - \theta) \rightsquigarrow \frac{W}{I(\theta)} \sim N\left(0, \frac{1}{I(\theta)}\right)$$

Assuming that $I(\theta)$ is a continuous function of θ , it follows that $I(\hat{\theta}_n) \xrightarrow{P} I(\theta)$. Now

$$\frac{\hat{\theta}_n - \theta}{\hat{\text{se}}} = \sqrt{n} I^{1/2}(\hat{\theta}_n)(\hat{\theta}_n - \theta) \tag{22}$$

$$= \left\{ \sqrt{n} I^{1/2}(\theta)(\hat{\theta}_n - \theta) \right\} \left\{ \frac{I(\hat{\theta}_n)}{I(\theta)} \right\}^{1/2} \tag{23}$$

The first term tends in distribution to $N(0, 1)$. The second term tends in probability to 1. The result follows from Theorem 6.5 part (e).

Outline of proof of Theorem 10.24. Write,

$$\hat{\tau} = g(\hat{\theta}) \approx g(\theta) + (\hat{\theta} - \theta)g'(\theta) = \tau + (\hat{\theta} - \theta)g'(\theta)$$

Thus,

$$\sqrt{n}(\hat{\tau} - \tau) \approx \sqrt{n}(\hat{\theta} - \theta)g'(\theta)$$

and hence

$$\frac{\sqrt{n}I^{1/2}(\theta)(\hat{\theta} - \theta)}{g'(\theta)} \approx \sqrt{n}I^{1/2}(\theta)(\hat{\theta} - \theta)$$

Theorem 10.18 tells us that the right hand side tends in distribution to $N(0, 1)$, hence

$$\frac{\sqrt{n}I^{1/2}(\theta)(\hat{\theta} - \theta)}{g'(\theta)} \rightsquigarrow N(0, 1)$$

or, in other words,

$$\hat{\tau} \approx N(\tau, \text{se}^2(\hat{\tau}_n))$$

where

$$\text{se}^2(\hat{\tau}_n) = \frac{(g'(\theta))^2}{nI(\theta)}$$

The result remains true if we substitute $\hat{\theta}$ for θ by Theorem 6.5 part (e).

10.12.2 Sufficiency

A **statistic** is a function $T(X^n)$ of the data. A sufficient statistic is a statistic that contains all of the information in the data.

Write $x^n \leftrightarrow y^n$ if $f(x^n; \theta) = cf(y^n; \theta)$ for some constant c that might depend on x^n and y^n but not θ . A statistic is **sufficient** if $T(x^n) \leftrightarrow T(y^n)$ implies that $x^n \leftrightarrow y^n$.

Notice that if $x^n \leftrightarrow y^n$ then the likelihood functions based on x^n and y^n have the same shape. Roughly speaking, a statistic is sufficient if we can calculate the likelihood function knowing only $T(X^n)$.

A statistic T is **minimally sufficient** if it is sufficient and it is a function of every other sufficient statistic.

Theorem 10.36. T is minimally sufficient if $T(x^n) = T(y^n)$ if and only if $x^n \leftrightarrow y^n$.

The usual definition of sufficiency is this: T is sufficient if the distribution of X^n given $T(X^n) = t$ does not depend on θ .

Theorem 10.40 (Factorization Theorem). T is sufficient if and only if there are functions $g(t, \theta)$ and $h(x)$ such that $f(x^n; \theta) = g(t(x^n); \theta)h(x^n)$.

Theorem 10.42 (Rao-Blackwell). Let $\hat{\theta}$ be an estimator and let T be a sufficient statistic. Define a new estimator by

$$\bar{\theta} = \mathbb{E}(\hat{\theta}|T)$$

Then, for every θ ,

$$R(\theta, \bar{\theta}) \leq R(\theta, \hat{\theta})$$

where $R(\theta, \hat{\theta}) = \mathbb{E}_{\theta}[(\theta - \hat{\theta})^2]$ denote the MSE of an estimator.

10.12.3 Exponential Families

We say that $\{f(x; \theta) : \theta \in \Theta\}$ is a **one-parameter exponential family** if there are functions $\eta(\theta)$, $B(\theta)$, $T(x)$ and $h(x)$ such that

$$f(x; \theta) = h(x)e^{\eta(\theta)T(x) - B(\theta)}$$

It is easy to see that $T(X)$ is sufficient. We call T the **natural sufficient statistic**.

We can rewrite an exponential family as

$$f(x; \eta) = h(x)e^{\eta T(x) - A(\eta)}$$

where $\eta = \eta(\theta)$ is called the **natural parameter** and

$$A(\eta) = \log \int h(x)e^{\eta T(x)} dx$$

Let X_1, \dots, X_n be iid from an exponential family. Then $f(x^n; \theta)$ is an exponential family:

$$f(x^n; \theta) = h_n(x^n)e^{\eta(\theta)T_n(x^n) - B_n(\theta)}$$

where $h_n(x^n) = \prod_i h(x_i)$, $T_n(x^n) = \sum_i T(x_i)$ and $B_n(\theta) = nB(\theta)$. This implies that $\sum_i T(X_i)$ is sufficient.

Theorem 10.47. Let X have an exponential family. Then,

$$\mathbb{E}(T(X)) = A'(\eta), \quad \mathbb{V}(T(X)) = A''(\eta)$$

If $\theta = (\theta_1, \dots, \theta_n)$ is a vector, then we say that $f(x; \theta)$ has exponential family form if

$$f(x; \theta) = h(x) \exp \left\{ \sum_{j=1}^k \eta_j(\theta) T_j(x) - B(\theta) \right\}$$

Again, $T = (T_1, \dots, T_k)$ is sufficient and n iid samples also has exponential form with sufficient statistic $(\sum_i T_1(X_i), \dots, \sum_i T_k(X_i))$.

10.13 Exercises

Exercise 10.13.1. Let $X_1, \dots, X_n \sim \text{Gamma}(\alpha, \beta)$. Find the method of moments estimator for α and β .

Solution.

The first two moments are:

$$\begin{aligned} \alpha_1 &= \mathbb{E}(X) = \frac{\alpha}{\beta} \\ \alpha_2 &= \mathbb{E}(X^2) = \mathbb{V}(X) + \mathbb{E}(X)^2 = \frac{\alpha}{\beta^2} + \frac{\alpha^2}{\beta^2} = \frac{\alpha(\alpha + 1)}{\beta^2} \end{aligned}$$

We have the sample moments:

$$\hat{\alpha}_1 = \frac{1}{n} \sum_{i=1}^n X_i \quad \hat{\alpha}_2 = \frac{1}{n} \sum_{i=1}^n X_i^2$$

Equating these we get:

$$\hat{\alpha}_1 = \frac{\hat{\alpha}_n}{\hat{\beta}_n} \quad \hat{\alpha}_2 = \frac{\hat{\alpha}_n(\hat{\alpha}_n + 1)}{\hat{\beta}_n^2}$$

Solving these we get the method of moments estimators for α and β :

$$\hat{\alpha}_n = \frac{\hat{\alpha}_1^2}{\hat{\alpha}_2 - \hat{\alpha}_1^2} \quad \hat{\beta}_n = \frac{\hat{\alpha}_1}{\hat{\alpha}_2 - \hat{\alpha}_1^2}$$

Exercise 10.13.2. Let $X_1, \dots, X_n \sim \text{Uniform}(a, b)$ where a, b are unknown parameters and $a < b$.

(a) Find the method of moments estimators for a and b .

(b) Find the MLE \hat{a} and \hat{b} .

(c) Let $\tau = \int x dF(x)$. Find the MLE of τ .

(d) Let $\hat{\tau}$ be the MLE from the previous item. Let $\tilde{\tau}$ be the nonparametric plug-in estimator of $\tau = \int x dF(x)$. Suppose that $a = 1$, $b = 3$ and $n = 10$. Find the MSE of $\hat{\tau}$ by simulation. Find the MSE of $\tilde{\tau}$ analytically. Compare.

Solution.

(a)

The first two moments are:

$$\alpha_1 = \mathbb{E}(X) = \frac{a+b}{2}$$

$$\alpha_2 = \mathbb{E}(X^2) = \mathbb{V}(X) + \mathbb{E}(X)^2 = \frac{(b-a)^2}{12} + \frac{(a+b)^2}{4} = \frac{a^2 + ab + b^2}{3}$$

We have the sample moments:

$$\hat{\alpha}_1 = \frac{1}{n} \sum_{i=1}^n X_i \quad \hat{\alpha}_2 = \frac{1}{n} \sum_{i=1}^n X_i^2$$

Equating these we get:

$$\hat{\alpha}_1 = \frac{\hat{a} + \hat{b}}{2} \quad \hat{\alpha}_2 = \frac{(\hat{b} - \hat{a})^2}{12} + \frac{(\hat{a} + \hat{b})^2}{4}$$

Solving these we get the method of moment estimators for a and b :

$$\hat{a} = \hat{\alpha}_1 - \sqrt{3}(\hat{\alpha}_1^2 - \hat{\alpha}_2) \quad \hat{b} = \hat{\alpha}_1 + \sqrt{3}(\hat{\alpha}_1^2 - \hat{\alpha}_2)$$

(b)

The probability density function for each X_i is

$$f(x; (a, b)) = \begin{cases} (b-a)^{-1} & \text{if } a \leq x \leq b \\ 0 & \text{otherwise} \end{cases}$$

The likelihood function is

$$\mathcal{L}_n(a, b) = \prod_{i=1}^n f(X_i; (a, b)) = \begin{cases} (b - a)^{-n} & \text{if } a \leq X_i \leq b \text{ for all } X_i \\ 0 & \text{otherwise} \end{cases}$$

The parameters that maximize the likelihood function make the $b - a$ as small as possible -- that is, we should pick the maximum a and the minimum b for which the likelihood function is non-zero. So the MLEs are:

$$\hat{a} = \min\{X_1, \dots, X_n\} \quad \hat{b} = \max\{X_1, \dots, X_n\}$$

(c)

$\tau = \int x dF(x) = \mathbb{E}(x) = (a + b)/2$, so since the MLE is equivariant, the MLE of τ is

$$\hat{\tau} = \frac{\hat{a} + \hat{b}}{2} = \frac{\min\{X_1, \dots, X_n\} + \max\{X_1, \dots, X_n\}}{2}$$

(d)

```
In [1]: import numpy as np

a = 1
b = 3
n = 10

X = np.random.uniform(low=a, high=b, size=n)

In [2]: tau_hat = (X.min() + X.max()) / 2

# Nonparametric bootstrap to find MSE of tau_hat
B = 10000
t_boot = np.empty(B)
for i in range(B):
    xx = np.random.choice(X, n, replace=True)
    t_boot[i] = (xx.min() + xx.max()) / 2

se = t_boot.std()
print("MSE for tau_hat: \t %.3f" % se)
```

```
MSE for tau_hat:      0.150
```

Analytically, we have:

$$\mathbb{V}(\tilde{\tau}) = \mathbb{E}(\tilde{\tau}^2) - (\mathbb{E}(\tilde{\tau}))^2 \quad (24)$$

$$= \frac{1}{n^2} \left(\mathbb{E} \left[\left(\sum_{i=1}^n X_i \right)^2 \right] - \left(\mathbb{E} \left[\sum_{i=1}^n X_i \right] \right)^2 \right) \quad (25)$$

$$= \frac{1}{n^2} \left(\mathbb{E} \left[\sum_{i=1}^n X_i^2 + \sum_{i=1}^n \sum_{j=1, j \neq i}^n X_i X_j \right] - \left(n \frac{a+b}{2} \right)^2 \right) \quad (26)$$

$$= \frac{1}{n^2} \left(\sum_{i=1}^n \mathbb{E}[X_i^2] + \sum_{i=1}^n \sum_{j=1, j \neq i}^n \mathbb{E}[X_i] \mathbb{E}[X_j] - \left(n \frac{a+b}{2} \right)^2 \right) \quad (27)$$

$$= \frac{1}{n^2} \left(n \frac{a^2 + ab + b^2}{3} + n(n-1) \left(\frac{a+b}{2} \right)^2 - n^2 \left(\frac{a+b}{2} \right)^2 \right) \quad (28)$$

$$= \frac{1}{n^2} \left(n \frac{a^2 + ab + b^2}{3} - n \left(\frac{a+b}{2} \right)^2 \right) \quad (29)$$

$$= \frac{1}{n} \left(\frac{a^2 + ab + b^2}{3} - \frac{a^2 + 2ab + b^2}{4} \right) \quad (30)$$

$$= \frac{1}{n} \frac{(b-a)^2}{12} \quad (31)$$

Therefore,

$$\text{se}(\tilde{\tau}) = \sqrt{\frac{1}{n} \frac{(b-a)^2}{12}}$$

```
In [3]: se_tau_tilde = np.sqrt((1/n) * ((b - a)**2 / 12))
        print("MSE for tau_tilde: \t %.3f" % se_tau_tilde)
```

MSE for tau_tilde: 0.183

Exercise 10.13.3. Let $X_1, \dots, X_n \sim N(\mu, \sigma^2)$. Let τ be the 0.95 percentile, i.e. $\mathbb{P}(X < \tau) = 0.95$.

(a) Find the MLE of τ .

(b) Find an expression for an approximate $1 - \alpha$ confidence interval for τ .

(c) Suppose the data are:

3.23	-2.50	1.88	-0.68	4.43	0.17
1.03	-0.07	-0.01	0.76	1.76	3.18
0.33	-0.31	0.30	-0.61	1.52	5.43
1.54	2.28	0.42	2.33	-1.03	4.00
0.39					

Find the MLE $\hat{\tau}$. Find the standard error using the delta method. Find the standard error using the parametric bootstrap.

Solution.

(a)

Let $Z \sim N(0, 1)$, so $(X - \mu)/\sigma \sim Z$. We have:

$$\mathbb{P}(X < \tau) = 0.95 \quad (32)$$

$$\mathbb{P}\left(\frac{X - \mu}{\sigma} < \frac{\tau - \mu}{\sigma}\right) = 0.95 \quad (33)$$

$$\mathbb{P}\left(Z < \frac{\tau - \mu}{\sigma}\right) = 0.95 \quad (34)$$

$$\frac{\tau - \mu}{\sigma} = z_{5\%} \quad (35)$$

$$\tau = \mu + z_{5\%}\sigma \quad (36)$$

Since the MLE is equivariant, $\hat{\tau} = \hat{\mu} + z_{5\%}\hat{\sigma}$, where $\hat{\mu}, \hat{\sigma}$ are the MLEs for the Normal distribution parameters:

$$\hat{\mu} = n^{-1} \sum_{i=1}^n X_i \quad \hat{\sigma} = \sqrt{n^{-1} \sum_{i=1}^n (X_i - \bar{X})^2}$$

(b)

Let's use the multiparameter delta method.

We have $\tau = g(\mu, \sigma) = \mu + z_{5\%}\sigma$, so

$$\nabla g = \begin{bmatrix} \partial g / \partial \mu \\ \partial g / \partial \sigma \end{bmatrix} = \begin{bmatrix} 1 \\ z_{5\%} \end{bmatrix}$$

The Fisher Information Matrix for the Normal process is

$$I_n(\mu, \sigma) = \begin{bmatrix} n/\sigma^2 & 0 \\ 0 & 2n/\sigma^2 \end{bmatrix}$$

then its inverse is

$$J_n = I_n^{-1}(\mu, \sigma) = \frac{1}{n} \begin{bmatrix} \sigma^2 & 0 \\ 0 & \sigma^2/2 \end{bmatrix}$$

and the standard error estimate for our new parameter variable is

$$\hat{\text{se}}(\hat{\tau}) = \sqrt{(\hat{\nabla} g)^T \hat{J}_n(\hat{\nabla} g)} = \hat{\sigma} \sqrt{n^{-1}(1 + z_{5\%}^2/2)}$$

A $1 - \alpha$ confidence interval for $\hat{\tau}$, then, is

$$C_n = \left(\hat{\mu} + \hat{\sigma} \left(z_{5\%} - z_{\alpha/2} \sqrt{n^{-1}(1 + z_{5\%}^2/2)} \right), \hat{\mu} + \hat{\sigma} \left(z_{5\%} + z_{\alpha/2} \sqrt{n^{-1}(1 + z_{5\%}^2/2)} \right) \right)$$

(c)

```
In [4]: import numpy as np
        from scipy.stats import norm

        z_05 = norm.ppf(0.95)
        z_025 = norm.ppf(0.975)
```

```
In [5]: X = np.array([
        3.23, -2.50, 1.88, -0.68, 4.43, 0.17,
        1.03, -0.07, -0.01, 0.76, 1.76, 3.18,
        0.33, -0.31, 0.30, -0.61, 1.52, 5.43,
        1.54, 2.28, 0.42, 2.33, -1.03, 4.00,
        0.39
    ])
```

```
In [6]: # Estimate the MLE tau_hat

n = len(X)
mu_hat = X.mean()
sigma_hat = X.std()
tau_hat = mu_hat + z_05 * sigma_hat

print("Estimated tau: %.3f" % tau_hat)
```

Estimated tau: 4.180

```
In [7]: # Confidence interval using delta method

se_tau_hat = sigma_hat * np.sqrt((1/n) * (1 + z_05 * z_05 / 2))
confidence_interval = (tau_hat - z_025 * se_tau_hat, tau_hat + z_025 * se_tau_hat)

print("Estimated tau (delta method, 95%% confidence interval): \t (%.3f, %.3f)" % confidence_interval)
```

Estimated tau (delta method, 95% confidence interval): (3.088, 5.273)

```
In [8]: # Confidence interval using parametric bootstrap

n = len(X)
mu_hat = X.mean()
sigma_hat = X.std()
tau_hat = mu_hat + z_05 * sigma_hat

B = 10000
t_boot = np.empty(B)
for i in range(B):
    xx = norm.rvs(loc=mu_hat, scale=sigma_hat, size=n)
    t_boot[i] = np.quantile(xx, 0.95)

se_tau_hat_bootstrap = t_boot.std()
confidence_interval = (tau_hat - z_025 * se_tau_hat_bootstrap, tau_hat + z_025 * se_tau_hat_bootstrap)

print("Estimated tau (parametric bootstrap, 95%% confidence interval): \t (%.3f, %.3f)" % confidence_interval)
```

Estimated tau (parametric bootstrap, 95% confidence interval): (2.887, 5.474)

Exercise 10.13.4 Let $X_1, \dots, X_n \sim \text{Uniform}(0, \theta)$. Show that the MLE is consistent.

Hint: Let $Y = \max\{X_1, \dots, X_n\}$. For any c , $\mathbb{P}(Y < c) = \mathbb{P}(X_1 < c, X_2 < c, \dots, X_n < c) = \mathbb{P}(X_1 < c)\mathbb{P}(X_2 < c) \dots \mathbb{P}(X_n < c)$.

Solution.

The probability density function is

$$f(x, \theta) = \mathbb{P}(Y < x) = \prod_{i=1}^n \mathbb{P}(X_i < x) = f_{\text{Uniform}(0, \theta)}(x)^n$$

The probability density function for the original distribution is

$$f_{\text{Uniform}(0, \theta)}(x) = \begin{cases} \theta^{-1} & \text{if } 0 \leq x \leq \theta \\ 0 & \text{otherwise} \end{cases}$$

so

$$f(x, \theta) = \begin{cases} \theta^{-n} & \text{if } 0 \leq x \leq \theta \\ 0 & \text{otherwise} \end{cases}$$

The likelihood is maximized when θ is as small as possible while keeping all samples within the first case, so $\hat{\theta}_n = \max\{X_1, \dots, X_n\}$.

For a given $\epsilon > 0$, we have

$$\mathbb{P}(\hat{\theta}_n < \theta - \epsilon) = \prod_{i=1}^n \mathbb{P}(X_i < \theta - \epsilon) = \left(1 - \frac{\epsilon}{\theta}\right)^n$$

which goes to 0 as $n \rightarrow \infty$, so $\lim_{n \rightarrow \infty} \hat{\theta}_n = \theta$, and thus the MLE is consistent.

Exercise 10.13.5. Let $X_1, \dots, X_n \sim \text{Poisson}(\lambda)$. Find the method of moments estimator, the maximum likelihood estimator, and the Fisher information $I(\lambda)$.

Solution.

The first moment is:

$$\mathbb{E}(X) = \lambda$$

We have the sample moment:

$$\hat{\alpha}_1 = \frac{1}{n} \sum_{i=1}^n X_i$$

Equating those, the method of moments estimator for $\hat{\lambda}$ is:

$$\hat{\lambda} = \hat{\alpha}_1 = \frac{1}{n} \sum_{i=1}^n X_i$$

The likelihood function is

$$\mathcal{L}_n(\lambda) = \prod_{i=1}^n f(X_i; \lambda) = \prod_{i=1}^n \frac{\lambda^{X_i} e^{-\lambda}}{(X_i)!}$$

so the log likelihood function is

$$\ell_n(\lambda) = \log \mathcal{L}_n(\lambda) = \sum_{i=1}^n (\log(\lambda^{X_i} e^{-\lambda}) - \log X_i!) = \sum_{i=1}^n (X_i \log \lambda - \lambda - \log X_i!) = -n\lambda + (\log \lambda) \sum_{i=1}^n X_i - \sum_{i=1}^n \log X_i!$$

To find the MLE, we differentiate this equation with respect to λ and equate it to 0:

$$\begin{aligned} \frac{\partial \ell_n(\lambda)}{\partial \lambda} &= 0 \\ -n + \frac{\sum_{i=1}^n X_i}{\lambda} &= 0 \\ \hat{\lambda} &= \frac{1}{n} \sum_{i=1}^n X_i \end{aligned}$$

The score function is:

$$s(X; \lambda) = \frac{\partial \log f(X; \lambda)}{\partial \lambda} = \frac{X}{\lambda} - 1$$

and the Fisher information is:

$$I_n(\lambda) = \sum_{i=1}^n \mathbb{V}(s(X_i; \lambda)) = \sum_{i=1}^n \mathbb{V}\left(\frac{X_i}{\lambda} - 1\right) = \frac{1}{\lambda^2} \sum_{i=1}^n \mathbb{V}(X_i) = \frac{n}{\lambda}$$

In particular, $I(\lambda) = I_1(\lambda) = 1/\lambda$.

Exercise 10.13.6. Let $X_1, \dots, X_n \sim N(\theta, 1)$. Define

$$Y_i = \begin{cases} 1 & \text{if } X_i > 0 \\ 0 & \text{if } X_i \leq 0 \end{cases}$$

Let $\psi = \mathbb{P}(Y_1 = 1)$.

(a) Find the maximum likelihood estimate $\hat{\psi}$ of ψ .

(b) Find an approximate 95% confidence interval for ψ .

(c) Define $\bar{\psi} = (1/n) \sum_i Y_i$. Show that $\bar{\psi}$ is a consistent estimator of ψ .

(d) Compute the asymptotic relative efficiency of $\bar{\psi}$ to $\hat{\psi}$. Hint: Use the delta method to get the standard error of the MLE. Then compute the standard error (i.e. the standard deviation) of $\bar{\psi}$.

(e) Suppose that the data are not really normal. Show that $\hat{\psi}$ is not consistent. What, if anything, does $\hat{\psi}$ converge to?

Solution.

Note that, from the definition, $Y_1, \dots, Y_n \sim \text{Bernoulli}(\Phi(\theta))$, where Φ is the CDF for the normal distribution. Let $p = \Phi(\theta)$.

(a) We have $\psi = \mathbb{P}(Y_1 = 1) = p$, so the MLE is $\hat{\psi} = \hat{p} = \Phi(\hat{\theta}) = \Phi(\bar{X})$, where $\bar{X} = n^{-1} \sum_{i=1}^n X_i$.

(b) Let $g(\theta) = \Phi(\theta)$. Then $g'(\theta) = \phi(\theta)$, where ϕ is the standard normal PDF. By the delta method, $\text{se}(\hat{\psi}) = |g'(\hat{\theta})| \text{se}(\hat{\theta}) = \phi(\bar{X}) n^{-1/2}$.

Then, an approximate 95% confidence interval is

$$C_n = \left(\Phi(\bar{X}) \left(1 - \frac{z_{2.5\%}}{\sqrt{n}} \right), \Phi(\bar{X}) \left(1 + \frac{z_{2.5\%}}{\sqrt{n}} \right) \right)$$

(c) $\bar{\psi}$ has mean p , so consistency follows from the law of large numbers.

(d) We have $\mathbb{V}(Y_1) = \psi(1 - \psi)$, since Y_1 follows a Bernoulli distribution, so $\mathbb{V}(\bar{\psi}) = \mathbb{V}(Y_1)/n = \psi(1 - \psi)/n$.

From (b), $\mathbb{V}\hat{\psi} = \phi(\theta)/n$.

Therefore, the asymptotic relative efficiency is

$$\frac{\psi(1 - \psi)}{\phi(\theta)} = \frac{\Phi(\theta)(1 - \Phi(\theta))}{\phi(\theta)}$$

(e) By the law of large numbers, we still have that \overline{X} converges to $\mathbb{E}(Y_1) = \mathbb{P}(Y_1 = 1) \cdot 1 + \mathbb{P}(Y_1 = 0) \cdot 0 = \mathbb{P}(Y_1 = 1) = 1 - F_X(0) = \mu_Y$. Then $\hat{\psi} = \Phi(\overline{X})$ converges to $\Phi(\mu_Y)$. But the true value of ψ is $\mathbb{P}(Y_1 = 1) = 1 - F_X(0)$.

But for an arbitrary distribution $1 - F_X(0) \neq \Phi(1 - F_X(0))$.

Exercise 10.13.7. (Comparing two treatments). n_1 people are given treatment 1 and n_2 people are given treatment 2. Let X_1 be the number of people on treatment 1 who respond favorably to the treatment and let X_2 be the number of people on treatment 2 who respond favorably. Assume that $X_1 \sim \text{Binomial}(n_1, p_1)$, $X_2 \sim \text{Binomial}(n_2, p_2)$. Let $\psi = p_1 - p_2$.

(a) Find the MLE of ψ .

(b) Find the Fisher Information Matrix $I(p_1, p_2)$.

(c) Use the multiparameter delta method to find the asymptotic standard error of $\hat{\psi}$.

(d) Suppose that $n_1 = n_2 = 200$, $X_1 = 160$ and $X_2 = 148$. Find $\hat{\psi}$. Find an approximate 90% confidence interval for ψ using (i) the delta method and (ii) the parametric bootstrap.

Solution.

(a) The MLE is equivariant, so

$$\hat{\psi} = \hat{p}_1 - \hat{p}_2 = \frac{X_1}{n_1} - \frac{X_2}{n_2}$$

(b)

The probability mass function is

$$f((x_1, x_2); \psi) = f_1(x_1; p_1) f_2(x_2; p_2) = \binom{n_1}{x_1} p_1^{x_1} (1 - p_1)^{n_1 - x_1} \binom{n_2}{x_2} p_2^{x_2} (1 - p_2)^{n_2 - x_2}$$

The log likelihood is

$$\ell_n = \log f((x_1, x_2); \psi) \tag{37}$$

$$= \sum_{i=1}^2 \log \binom{n_i}{x_i} + x_i \log p_i + (n_i - x_i) \log(1 - p_i) \tag{38}$$

Calculating the partial derivatives and their expectations,

$$H_{11} = \frac{\partial^2 \ell_n}{\partial p_1^2} = \frac{\partial}{\partial p_1} \left(\frac{x_1}{p_1} - \frac{n_1 - x_1}{1 - p_1} \right) = -\frac{x_1}{p_1^2} - \frac{n_1 - x_1}{(1 - p_1)^2} \quad (39)$$

$$\mathbb{E}[H_{11}] = -\frac{\mathbb{E}[x_1]}{p_1^2} - \frac{\mathbb{E}[n - x_1]}{(1 - p_1)^2} = -\frac{n_1/p_1}{p_1^2} - \frac{n_1/(1 - p_1)}{(1 - p_1)^2} = -\frac{n_1}{p_1} - \frac{n_1}{1 - p_1} = -\frac{n_1}{p_1(1 - p_1)} \quad (40)$$

$$H_{22} = -\frac{x_2}{p_2^2} - \frac{n_2 - x_2}{(1 - p_2)^2} \quad (41)$$

$$\mathbb{E}[H_{22}] = -\frac{n_2}{p_2(1 - p_2)} \quad (42)$$

$$H_{12} = \frac{\partial^2 \ell_n}{\partial p_1 \partial p_2} = 0$$

$$H_{21} = 0$$

So the Fisher Information Matrix is:

$$I(p_1, p_2) = \begin{bmatrix} \frac{n_1}{p_1(1-p_1)} & 0 \\ 0 & \frac{n_2}{p_2(1-p_2)} \end{bmatrix}$$

(c) Using the multiparameter delta method, $g(\psi) = p_1 - p_2$, so

$$\nabla g = \begin{bmatrix} \partial g / \partial p_1 \\ \partial g / \partial p_2 \end{bmatrix} = \begin{bmatrix} 1 \\ -1 \end{bmatrix}$$

The inverse of the Fisher Information Matrix is

$$J(p_1, p_2) = I(p_1, p_2)^{-1} = \begin{bmatrix} \frac{p_1(1-p_1)}{n_1} & 0 \\ 0 & \frac{p_2(1-p_2)}{n_2} \end{bmatrix}$$

Then the asymptotic standard error of $\hat{\psi}$ is:

$$\hat{\text{se}}(\hat{\psi}) = \sqrt{(\hat{\nabla}g)^T \hat{J}_n(\hat{\nabla}g)} = \sqrt{\frac{p_1(1-p_1)}{n_1} + \frac{p_2(1-p_2)}{n_2}}$$

(d)

```
In [9]: import numpy as np
        from scipy.stats import norm, binom
```

```
n = 200
X1 = 160
X2 = 148
```

```
In [10]: p1_hat = X1 / n
          p2_hat = X2 / n
          psi_hat = p1_hat - p2_hat

          print("Estimated psi: \t %.3f" % psi_hat)
```

Estimated psi: 0.060

```
In [11]: # Confidence using delta method

          z = norm.ppf(.95)

          se_delta = np.sqrt(p1_hat * (1 - p1_hat) / n + p2_hat * (1 - p2_hat) / n)
          confidence_delta = (psi_hat - z * se_delta, psi_hat + z * se_delta)

          print("90%% confidence interval (delta method): \t %.3f, %.3f" % confidence_delta)
```

90% confidence interval (delta method): -0.009, 0.129

```
In [12]: # Confidence using parametric bootstrap

          B = 1000
          xx1 = binom.rvs(n, p1_hat, size=B)
          xx2 = binom.rvs(n, p2_hat, size=B)
          t_boot = xx1 / n - xx2 / n

          se_bootstrap = t_boot.std()
          confidence_delta = (psi_hat - z * se_bootstrap, psi_hat + z * se_bootstrap)

          print("90%% confidence interval (parametric bootstrap): \t %.3f, %.3f" % confidence_delta)
```

90% confidence interval (parametric bootstrap): -0.010, 0.130

Exercise 10.13.8. Find the Fisher information matrix for Example 10.29:

Let $X_1, \dots, X_n \sim N(\mu, \sigma^2)$.

Solution The log likelihood is:

$$\ell_n = \sum_i \log f(x; (\mu, \sigma)) = n \left[\log \left(\frac{1}{\sigma \sqrt{2\pi}} \right) + \left(-\frac{1}{2} \left(\frac{x - \mu}{\sigma} \right)^2 \right) \right]$$

From this,

$$H_{11} = \frac{\partial^2 \ell_n}{\partial \mu^2} = -\frac{n}{\sigma^2}$$

$$H_{22} = \frac{\partial^2 \ell_n}{\partial \sigma^2} = -\frac{n}{\sigma^2} - \frac{n}{\sigma^2} = -\frac{2n}{\sigma^2}$$

$$H_{12} = H_{21} = \frac{\partial^2 \ell_n}{\partial \mu \partial \sigma} = 0$$

So the Fisher Information Matrix is

$$I(\mu, \sigma) = - \begin{bmatrix} \mathbb{E}[H_{11}] & \mathbb{E}[H_{12}] \\ \mathbb{E}[H_{21}] & \mathbb{E}[H_{22}] \end{bmatrix} = \begin{bmatrix} \frac{n}{\sigma^2} & 0 \\ 0 & \frac{2n}{\sigma^2} \end{bmatrix}$$

Exercises 10.13.9 and 10.13.10. See final exercises from chapter 9.