

# 15. Multivariate Models

Review of notation from linear algebra:

- If  $x$  and  $y$  are vectors, then  $x^T y = \sum_j x_j y_j$ .
- If  $A$  is a matrix then  $\det(A)$  denotes the determinant of  $A$ ,  $A^T$  denotes the transpose of  $A$ , and  $A^{-1}$  denotes the inverse of  $A$  (if the inverse exists).
- The trace of a square matrix  $A$ , denoted by  $\text{tr}(A)$ , is the sum of its diagonal elements.
- The trace satisfies  $\text{tr}(AB) = \text{tr}(BA)$  and  $\text{tr}(A + B) = \text{tr}(A) + \text{tr}(B)$ .
- The trace satisfies  $\text{tr}(a) = a$  if  $a$  is a scalar.
- A matrix  $\Sigma$  is positive definite if  $x^T \Sigma x > 0$  for all non-zero vectors  $x$ .
- If a matrix  $\Sigma$  is symmetric and positive definite, there exists a matrix  $\Sigma^{1/2}$ , called the square root of  $\Sigma$ , with the following properties:
  - $\Sigma^{1/2}$  is symmetric
  - $\Sigma = \Sigma^{1/2} \Sigma^{1/2}$
  - $\Sigma^{1/2} \Sigma^{-1/2} = \Sigma^{-1/2} \Sigma^{1/2} = I$  where  $\Sigma^{-1/2} = (\Sigma^{1/2})^{-1}$ .

## 15.1 Random Vectors

Multivariate models involve a random vector  $X$  of the form

$$X = \begin{pmatrix} X_1 \\ \vdots \\ X_k \end{pmatrix}$$

The mean of a random vector  $X$  is defined by

$$\mu = \begin{pmatrix} \mu_1 \\ \vdots \\ \mu_k \end{pmatrix} = \begin{pmatrix} \mathbb{E}(X_1) \\ \vdots \\ \mathbb{E}(X_k) \end{pmatrix}$$

The covariance matrix  $\Sigma$  is defined to be

$$\Sigma = \mathbb{V}(X) = \begin{pmatrix} \mathbb{V}(X_1) & \text{Cov}(X_1, X_2) & \cdots & \text{Cov}(X_1, X_k) \\ \text{Cov}(X_2, X_1) & \mathbb{V}(X_2) & \cdots & \text{Cov}(X_2, X_k) \\ \vdots & \vdots & \ddots & \vdots \\ \text{Cov}(X_k, X_1) & \text{Cov}(X_k, X_2) & \cdots & \mathbb{V}(X_k) \end{pmatrix}$$

This is also called the variance matrix or the variance-covariance matrix.

**Theorem 15.1.** Let  $a$  be a vector of length  $k$  and let  $X$  be a random vector of the same length with mean  $\mu$  and variance  $\Sigma$ . Then  $\mathbb{E}(a^T X) = a^T \mu$  and  $\mathbb{V}(a^T X) = a^T \Sigma a$ . If  $A$  is a matrix with  $k$  columns then  $\mathbb{E}(AX) = A\mu$  and  $\mathbb{V}(AX) = A\Sigma A^T$ .

Now suppose we have a random sample of  $n$  vectors:

$$\begin{pmatrix} X_{11} \\ X_{21} \\ \vdots \\ X_{k1} \end{pmatrix}, \begin{pmatrix} X_{21} \\ X_{22} \\ \vdots \\ X_{k2} \end{pmatrix}, \dots, \begin{pmatrix} X_{1n} \\ X_{2n} \\ \vdots \\ X_{kn} \end{pmatrix}$$

The sample mean  $\bar{X}$  is a vector defined by

$$\bar{X} = \begin{pmatrix} \bar{X}_1 \\ \vdots \\ \bar{X}_k \end{pmatrix}$$

where  $\bar{X}_i = n^{-1} \sum_{j=1}^n X_{ij}$ . The sample variance matrix is

$$S = \begin{pmatrix} s_{11} & s_{12} & \cdots & s_{1k} \\ s_{12} & s_{22} & \cdots & s_{2k} \\ \vdots & \vdots & \ddots & \vdots \\ s_{1k} & s_{2k} & \cdots & s_{kk} \end{pmatrix}$$

where

$$s_{ab} = \frac{1}{n-1} \sum_{j=1}^n (X_{aj} - \bar{X}_a)(X_{bj} - \bar{X}_b)$$

It follows that  $\mathbb{E}(\overline{X}) = \mu$  and  $\mathbb{E}(S) = \Sigma$ .

## 15.2 Estimating the Correlation

Consider  $n$  data points from a bivariate distribution

$$\begin{pmatrix} X_{11} \\ X_{21} \end{pmatrix}, \begin{pmatrix} X_{12} \\ X_{22} \end{pmatrix}, \dots, \begin{pmatrix} X_{1n} \\ X_{2n} \end{pmatrix}$$

Recall that the correlation between  $X_1$  and  $X_2$  is

$$\rho = \frac{\mathbb{E}((X_1 - \mu)(X_2 - \mu_2))}{\sigma_1 \sigma_2}$$

The sample correlation (the plug-in estimator) is

$$\hat{\rho} = \frac{\sum_{i=1}^n (X_{1i} - \overline{X}_1)(X_{2i} - \overline{X}_2)}{s_1 s_2}$$

We can construct a confidence interval for  $\rho$  by applying the delta method as usual. However, it turns out that we get a more accurate confidence interval by first constructing a confidence interval for a function  $\theta = f(\rho)$  and then applying the inverse function  $f^{-1}$ . The method, due to Fisher, is as follows. Define

$$f(r) = \frac{1}{2}(\log(1+r) - \log(1-r))$$

and let  $\theta = f(\rho)$ . The inverse of  $f$  is

$$g(z) \equiv f^{-1}(z) = \frac{e^{2z} - 1}{e^{2z} + 1}$$

Now do the following steps:

### Approximate Confidence Interval for the Correlation

1. Compute

$$\hat{\theta} = f(\hat{\rho}) = \frac{1}{2}(\log(1 + \hat{\rho}) - \log(1 - \hat{\rho}))$$

1. Compute the approximate standard error of  $\hat{\theta}$  which can be shown to be

$$\hat{\text{se}}(\hat{\theta}) = \frac{1}{\sqrt{n-3}}$$

1. An approximate  $1 - \alpha$  confidence interval for  $\theta = f(\rho)$  is

$$(a, b) \equiv \left( \hat{\theta} - \frac{z_{\alpha/2}}{\sqrt{n-3}}, \hat{\theta} + \frac{z_{\alpha/2}}{\sqrt{n-3}} \right)$$

1. Apply the inverse transformation  $f^{-1}(z)$  to get a confidence interval for  $\rho$ :

$$\left( \frac{e^{2a} - 1}{e^{2a} + 1}, \frac{e^{2b} - 1}{e^{2b} + 1} \right)$$

## 15.3 Multinomial

Review of Multinomial distribution: consider drawing a ball from an urn that has  $n$  balls of  $k$  colors. Let  $p = (p_1, \dots, p_k)$  where  $p_j \geq 0$  are the probabilities of drawing (with replacement) a ball of each color;  $\sum_j p_j = 1$ . Draw  $n$  times and let  $X = (X_1, \dots, X_n)$  where  $X_j$  is the number of times that color  $j$  appeared; so  $\sum_k X_k = n$ . We say  $X$  has a Multinomial( $n, p$ ) distribution. The probability function is

$$f(x; p) = \binom{n}{x_1 \dots x_k} p_1^{x_1} \dots p_k^{x_k}$$

where

$$\binom{n}{x_1 \dots x_k} = \frac{n!}{x_1! \dots x_k!}$$

**Theorem 15.2.** Let  $X \sim \text{Multinomial}(n, p)$ . Then the marginal distribution of  $X_j$  is  $X_j \sim \text{Binomial}(n, p_j)$ . The mean and variance of  $X$  are

$$\mathbb{E}(X) = \begin{pmatrix} np_1 \\ \vdots \\ np_k \end{pmatrix}$$

and

$$\mathbb{V}(X) = \begin{pmatrix} np_1(1-p_1) & -np_1p_2 & \cdots & -np_1p_k \\ -np_1p_2 & np_2(1-p_2) & \cdots & -np_2p_k \\ \vdots & \vdots & \ddots & \vdots \\ -np_1p_k & -np_2p_k & \cdots & np_k(1-p_k) \end{pmatrix}$$

**Proof.** That  $X_j \sim \text{Binomial}(n, p_j)$  follows easily. Hence  $\mathbb{E}(X_j) = np_j$  and  $\mathbb{V}(X_j) = np_j(1-p_j)$ .

To compute  $\text{Cov}(X_i, X_j)$ , notice that  $X_i + X_j \sim \text{Binomial}(n, p_i + p_j)$ , so  $\mathbb{V}(X_i + X_j) = n(p_i + p_j)(1 - p_i - p_j)$ . On the other hand, decomposing the sum of the random variables on the variance,

$$\mathbb{V}(X_i + X_j) = \mathbb{V}(X_i) + \mathbb{V}(X_j) + 2\text{Cov}(X_i, X_j) \quad (1)$$

$$= np_i(1-p_i) + np_j(1-p_j) + 2\text{Cov}(X_i, X_j) \quad (2)$$

Equating both expressions and isolating the covariance we get  $\text{Cov}(X_i, X_j) = -np_ip_j$ .

**Theorem 15.3.** The maximum likelihood estimator of  $p$  is

$$\hat{p} = \begin{pmatrix} \hat{p}_1 \\ \vdots \\ \hat{p}_k \end{pmatrix} = \begin{pmatrix} \frac{X_1}{n} \\ \vdots \\ \frac{X_k}{n} \end{pmatrix} = \frac{X}{n}$$

**Proof.** The log-likelihood (ignoring a constant) is  $\ell(p) = \sum_j X_j \log p_j$ . When we maximize it we need to be careful to enforce the constraint that  $\sum_j p_j = 1$ . Using Lagrange multipliers, instead we maximize

$$A(p) = \sum_{j=1}^k X_j \log p_j + \lambda \left( \sum_{j=1}^k p_j - 1 \right)$$

But

$$\frac{\partial A(p)}{\partial p_j} = \frac{X_j}{p_j} + \lambda$$

Setting it to zero we get  $\hat{p}_j = -X_j/\lambda$ . Since  $\sum_j \hat{p}_j = 1$  we get  $\lambda = -n$  and so  $\hat{p}_j = X_j/n$ , which is our result.

Next we want the variance of the MLE. The direct approach is to compute the variance matrix of  $\hat{p}$  directly:  $\mathbb{V}(\hat{p}) = \mathbb{V}(X/n) = n^{-2}\mathbb{V}(X)$ , so

$$\mathbb{V}(\hat{p}) = \frac{1}{n} \begin{pmatrix} p_1(1-p_1) & -p_1p_2 & \cdots & -p_1p_k \\ -p_1p_2 & p_2(1-p_2) & \cdots & -p_2p_k \\ \vdots & \vdots & \ddots & \vdots \\ -p_1p_k & -p_2p_k & \cdots & p_k(1-p_k) \end{pmatrix}$$

## 15.4 Multivariate Normal

Let's recall the definition of the multivariate normal. Let

$$Z = \begin{pmatrix} Z_1 \\ \vdots \\ Z_k \end{pmatrix}$$

where  $Z_1, \dots, Z_k \sim N(0, 1)$  are independent. The density of  $Z$  is

$$f(z) = \frac{1}{(2\pi)^{k/2}} \exp \left\{ -\frac{1}{2} \sum_{j=1}^k z_j^2 \right\} = \frac{1}{(2\pi)^{k/2}} \exp \left\{ -\frac{1}{2} z^T z \right\}$$

The variance matrix of  $Z$  is the identity matrix  $I$ . We write  $Z \sim N(0, I)$  where it is understood that  $0$  is a vector of  $k$  zeroes. We say  $Z$  has a standard multivariate Normal distribution.

More generally, a vector  $X$  has a multivariate Normal distribution, denoted by  $X \sim N(\mu, \Sigma)$ , if its density is

$$f(x; \mu, \Sigma) = \frac{1}{(2\pi)^{k/2} \det(\Sigma)^{1/2}} \exp \left\{ -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right\}$$

where  $\mu$  is a vector of length  $k$  and  $\Sigma$  is a  $k \times k$  symmetric, positive definite matrix. Then  $\mathbb{E}(X) = \mu$  and  $\mathbb{V}(X) = \Sigma$ . Setting  $\mu = 0$  and  $\Sigma = I$  gives back the standard Normal.

**Theorem 15.4.** The following properties hold:

1. If  $Z \sim N(0, 1)$  and  $X = \mu + \Sigma^{1/2}Z$  then  $X \sim N(\mu, \Sigma)$ .
2. If  $X \sim N(\mu, \Sigma)$ , then  $\Sigma^{-1/2}(X - \mu) \sim N(0, 1)$ .
3. If  $X \sim N(\mu, \Sigma)$  and  $a$  is a vector with the same length as  $X$ , then  $a^T X \sim N(a^T \mu, a^T \Sigma a)$ .

4. Let  $V = (X - \mu)^T \Sigma^{-1} (X - \mu)$ . Then  $V \sim \xi_k^2$ .

Suppose we partition a random Normal vector  $X$  into two parts  $X = (X_a, X_b)$ . We can similarly partition the mean  $\mu = (\mu_a, \mu_b)$  and the variance

$$\Sigma = \begin{pmatrix} \Sigma_{aa} & \Sigma_{ab} \\ \Sigma_{ba} & \Sigma_{bb} \end{pmatrix}$$

**Theorem 15.5.** Let  $X \sim N(\mu, \Sigma)$ . Then:

1. The marginal distribution of  $X_a$  is  $X_a \sim N(\mu_a, \Sigma_{aa})$ .
2. The conditional distribution of  $X_b$  given  $X_a = x_a$  is

$$X_b | X_a = x_a \sim N(\mu(x_a), \Sigma(x_a))$$

where

$$\mu(x_a) = \mu_b + \Sigma_{ba} \Sigma_{aa}^{-1} (x_a - \mu_a) \quad (3)$$

$$\Sigma(x_a) = \Sigma_{bb} - \Sigma_{ba} \Sigma_{aa}^{-1} \Sigma_{ab} \quad (4)$$

**Theorem 15.6.** Given a random sample of size  $n$  from a  $N(\mu, \Sigma)$ , the log-likelihood (up to a constant not depending on  $\mu$  or  $\Sigma$ ) is given by

$$\ell(\mu, \Sigma) = -\frac{n}{2} (\bar{X} - \mu)^T \Sigma^{-1} (\bar{X} - \mu) - \frac{n}{2} \text{tr}(\Sigma^{-1} S) - \frac{n}{2} \log \det(\Sigma)$$

The MLE is

$$\hat{\mu} = \bar{X} \quad \text{and} \quad \hat{\Sigma} = \left( \frac{n-1}{n} \right) S$$

## 15.5 Appendix

**Proof of Theorem 15.6.** Let the  $i$ -th random vector be  $X^i$ . The log-likelihood is

$$\ell(\mu, \Sigma) = \sum_{i=1}^k f(X^i; \mu, \Sigma) = -\frac{kn}{2} \log(2\pi) - \frac{n}{2} \log \det(\Sigma) - \frac{1}{2} \sum_{i=1}^k (X^i - \mu)^T \Sigma^{-1} (X^i - \mu)$$

Now,

$$\sum_{i=1}^k (X^i - \mu)^T \Sigma^{-1} (X^i - \mu) = \sum_{i=1}^k \left[ (X^i - \bar{X}) + (\bar{X} - \mu) \right]^T \Sigma^{-1} \left[ (X^i - \bar{X}) + (\bar{X} - \mu) \right] \quad (5)$$

$$= \sum_{i=1}^k [(X^i - \bar{X})^T \Sigma^{-1} (X^i - \bar{X})] + n(\bar{X} - \mu)^T \Sigma^{-1} (\bar{X} - \mu) \quad (6)$$

since  $\sum_i (X^i - \bar{X}) \Sigma^{-1} (\bar{X} - \mu) = 0$ . Also,  $(X^i - \mu)^T \Sigma^{-1} (X^i - \mu)$  is a scalar, so

$$\sum_{i=1}^k (X^i - \mu)^T \Sigma^{-1} (X^i - \mu) = \sum_{i=1}^k \text{tr} [(X^i - \mu)^T \Sigma^{-1} (X^i - \mu)] \quad (7)$$

$$= \sum_{i=1}^k \text{tr} [\Sigma^{-1} (X^i - \mu) (X^i - \mu)^T] \quad (8)$$

$$= \text{tr} \left[ \Sigma^{-1} \sum_{i=1}^k (X^i - \mu) (X^i - \mu)^T \right] \quad (9)$$

$$= n \text{tr} [\Sigma^{-1} S] \quad (10)$$

so the conclusion follows.

## 15.6 Exercises

**Exercise 15.6.1.** Prove Theorem 15.1.

Let  $a$  be a vector of length  $k$  and let  $X$  be a random vector of the same length with mean  $\mu$  and variance  $\Sigma$ . Then  $\mathbb{E}(a^T X) = a^T \mu$  and  $\mathbb{V}(a^T X) = a^T \Sigma a$ . If  $A$  is a matrix with  $k$  columns then  $\mathbb{E}(AX) = A\mu$  and  $\mathbb{V}(AX) = A\Sigma A^T$ .

**Solution.**

For the vector version of the theorem, we have:

$$\mathbb{E}(a^T X) = \mathbb{E} \left( \sum_{i=1}^k a_i X_i \right) = \sum_{i=1}^k a_i \mathbb{E}(X_i) = \sum_{i=1}^k a_i \mu_i = a^T \mu$$

$$\mathbb{V}(a^T X) = \mathbb{V} \left( \sum_{i=1}^k a_i X_i \right) = \sum_{i=1}^k \sum_{j=1}^k a_i a_j \text{Cov}(X_i, X_j) = \sum_{i=1}^k a_i \left( \sum_{j=1}^k \text{Cov}(X_i, X_j) a_j \right) = \sum_{i=1}^k a_i (\Sigma a)_i = a^T \Sigma a$$



For the matrix version of the theorem, consider the  $r$  rows of  $A$  as vectors, separately,  $a^1, \dots, a^k$ :

$$A = \begin{pmatrix} \dots & a^1 & \dots \\ \dots & a^2 & \dots \\ \vdots & \vdots & \vdots \\ \dots & a^r & \dots \end{pmatrix}$$

Then,

$$\mathbb{E}(AX) = \begin{pmatrix} \mathbb{E}(a^1 X) \\ \mathbb{E}(a^2 X) \\ \vdots \\ \mathbb{E}(a^r X) \end{pmatrix} = \begin{pmatrix} a^1 \mu \\ a^2 \mu \\ \vdots \\ a^r \mu \end{pmatrix} = A\mu$$

Finally, looking at the  $i$ -th term of  $AX$ ,

$$(AX)_i = \sum_{s=1}^k a_{is} X_s = a^i X$$

so, by the vector version of the theorem,  $\mathbb{V}((AX)_i) = (a^i)^T \Sigma a^i$ . Applying this to every element:

$$\mathbb{V}(AX) = \begin{pmatrix} \mathbb{V}((AX)_1) \\ \mathbb{V}((AX)_2) \\ \vdots \\ \mathbb{V}((AX)_r) \end{pmatrix} = \begin{pmatrix} \mathbb{V}(a^1 X) \\ \mathbb{V}(a^2 X) \\ \vdots \\ \mathbb{V}(a^r X) \end{pmatrix} = \begin{pmatrix} (a^1)^T \Sigma a^1 \\ (a^2)^T \Sigma a^2 \\ \vdots \\ (a^r)^T \Sigma a^r \end{pmatrix} = A \Sigma A^T$$

**Exercise 15.6.2.** Find the Fisher information matrix for the MLE of a Multinomial.

**Solution.**

The probability mass function for a Multinomial distribution is:

$$f(X; p) = \binom{n}{X_1 \dots X_k} p_1^{X_1} \dots p_k^{X_k} = \frac{n!}{X_1! \dots X_k!} p_1^{X_1} \dots p_k^{X_k}$$

so the log-likelihood (ignoring a constant) is

$$\ell_n(p) = \sum_{i=1}^k X_i \log p_i$$

The partial derivatives are:

$$H_{ii} = \frac{\partial^2 \ell_n(p)}{\partial^2 p_i} = -\frac{X_i}{p_i^2} \quad (11)$$

$$H_{ij} = \frac{\partial^2 \ell_n(p)}{\partial p_i \partial p_j} = 0 \text{ for } i \neq j \quad (12)$$

so  $\mathbb{E}(H_{ii}) = -n/p_i$ , and the Fisher Information Matrix is:

$$I_n(p) = n \begin{pmatrix} \frac{1}{p_1} & 0 & \cdots & 0 \\ 0 & \frac{1}{p_2} & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \frac{1}{p_k} \end{pmatrix}$$

Note, however, that the variance is *not* the inverse matrix of  $I_n(p)$ , and further, that,  $\mathbb{V}(X)$  is not invertible.

**Exercise 15.6.3.** Prove Theorem 15.5.

Let  $X \sim N(\mu, \Sigma)$ . Then:

1. The marginal distribution of  $X_a$  is  $X_a \sim N(\mu_a, \Sigma_{aa})$ .
2. The conditional distribution of  $X_b$  given  $X_a = x_a$  is

$$X_b | X_a = x_a \sim N(\mu(x_a), \Sigma(x_a))$$

where

$$\mu(x_a) = \mu_b + \Sigma_{ba} \Sigma_{aa}^{-1} (x_a - \mu_a) \quad (13)$$

$$\Sigma(x_a) = \Sigma_{bb} - \Sigma_{ba} \Sigma_{aa}^{-1} \Sigma_{ab} \quad (14)$$

**Solution.**

The marginal distribution result is immediate: for any given sample drawn from the distribution, collect only the first  $k$  dimensions of the sample vector, where  $k$  is the number of dimensions of  $X_a$ . The resulting distribution will be multivariate normal, with mean and variance given by getting the first  $k$  dimensions of  $\mu$  and  $\Sigma$ .

For the conditional distribution result, let  $A = -\Sigma_{ba}\Sigma_{aa}^{-1}$  and  $z = x_b + Ax_a$ . We have:

$$\text{Cov}(z, x_a) = \text{Cov}(x_b, x_a) + \text{Cov}(Ax_a, x_a) \quad (15)$$

$$= \Sigma_{ba} + A\Sigma_{aa} \quad (16)$$

$$= \Sigma_{ba} - \Sigma_{ba}\Sigma_{aa}^{-1}\Sigma_{aa} \quad (17)$$

$$= 0 \quad (18)$$

so  $z$  and  $x_a$  are uncorrelated (and since they are jointly normal, they are also independent). We then have:

$$\mathbb{E}(x_b|x_a) = \mathbb{E}(z - Ax_a|x_a) \quad (19)$$

$$= \mathbb{E}(z|x_a) - \mathbb{E}(Ax_a|x_a) \quad (20)$$

$$= \mathbb{E}(z) - Ax_a \quad (21)$$

$$= \mu_b + A\mu_a - Ax_a \quad (22)$$

$$= \mu_b + \Sigma_{ba}\Sigma_{aa}^{-1}(x_a - \mu_a) \quad (23)$$

For the covariance matrix,

$$\mathbb{V}(x_b|x_a) = \mathbb{V}(z - Ax_a|x_a) \quad (24)$$

$$= \mathbb{V}(z|x_a) - \mathbb{V}(Ax_a|x_a) - A\text{Cov}(z, -x_a) - \text{Cov}(z, -x_a)A^T \quad (25)$$

$$= \mathbb{V}(z|x_a) - 0 - A \cdot 0 - 0 \cdot A \quad (26)$$

$$= \mathbb{V}(z) \quad (27)$$

$$= \mathbb{V}(x_b + Ax_a) \quad (28)$$

$$= \mathbb{V}(x_b) + A\mathbb{V}(x_a)A^T + A\text{Cov}(x_a, x_b) + \text{Cov}(x_b, x_a)A^T \quad (29)$$

$$= \Sigma_{bb} + (-\Sigma_{ba}\Sigma_{aa}^{-1})\Sigma_{aa}(-\Sigma_{ba}\Sigma_{aa}^{-1})^T + (-\Sigma_{ba}\Sigma_{aa}^{-1})\Sigma_{ab} + \Sigma_{ba}(-\Sigma_{ba}\Sigma_{aa}^{-1})^T \quad (30)$$

$$= \Sigma_{bb} + \Sigma_{ba}\Sigma_{aa}^{-1}\Sigma_{aa}\Sigma_{aa}^{-1}\Sigma_{ab} - 2\Sigma_{ba}\Sigma_{aa}^{-1}\Sigma_{ab} \quad (31)$$

$$= \Sigma_{bb} - \Sigma_{ba}\Sigma_{aa}^{-1}\Sigma_{ab} \quad (32)$$

Reference: Macro (<https://stats.stackexchange.com/users/4856/macro>), Deriving the conditional distributions of a multivariate normal distribution, URL (version: 2015-06-18): <https://stats.stackexchange.com/q/30600>

**Exercise 15.6.4 (Computer Experiment).** Write a function to generate `nsim` observations from a `Multinomial( $n, p$ )` distribution.

**Solution.** Let's use the combinatoric interpretation of the distribution: Drawing  $n$  times, with replacement, from an urn with different ball colors, where the probability of obtaining balls of color  $i$  is  $p_i$ .

```
In [1]: import numpy as np

def multinomial_observations(n, p, nsim=1):
    cumulative_probabilities = np.cumsum(p)

    # Ensure probabilities add up to 1 (approximately)
    assert abs(cumulative_probabilities[-1] - 1) < 1e-8, "Probabilities should add up to 1"

    def get_observation():
        counts = np.zeros(cumulative_probabilities.size).astype(int)
        rvs = np.random.uniform(size=n)
        for i in range(n):
            counts[np.argmax(rvs[i] > cumulative_probabilities)] += 1


        return counts

    return np.array([get_observation() for _ in range(nsim)])
```

```
In [2]: # Sample usage
import matplotlib.pyplot as plt
%matplotlib inline

obs = multinomial_observations(n=500, p=[0.5, 0.3, 0.2], nsim=1000)

plt.figure(figsize=(12, 8))
plt.hist(obs[:, 0], density=True, bins=25, histtype='step', color='red', label='p = 0.5')
plt.hist(obs[:, 1], density=True, bins=25, histtype='step', color='green', label='p = 0.3')
plt.hist(obs[:, 2], density=True, bins=25, histtype='step', color='blue', label='p = 0.2')
plt.legend(loc='upper right')
plt.show()
```



**Exercise 15.6.5 (Computer Experiment).** Write a function to generate `nsim` observations from a Multivariate normal with given mean  $\mu$  and covariance matrix  $\Sigma$ .

**Solution.** Let's construct our samples based on samples of a standard multivariate normal  $Z \sim N(0, I)$ , by making  $X = \mu + \Sigma^{1/2}Z$ .

```
In [3]: import numpy as np

def multivariate_normal_observations(mu, sigma, nsim=1):
    mu_array = np.array(mu)
    sigma_array = np.array(sigma)

    assert len(mu_array.shape) == 1, "mu should be a vector"
    k = mu_array.shape[0]
```

```

assert sigma_array.shape == (k, k), "sigma should be a square matrix with same length as mu"

# Do the eigenvalue decomposition, then get  $U D^{1/2}$  as  $\Sigma^{1/2}$ 
U, D, V = np.linalg.svd(sigma_array)
sigma_sqrt = U @ np.diag(np.sqrt(D))

# Let's write our own random normal generator for fun, rather than use np.random.normal
# Strategy: Use Box-Muller to transform two random uniform variables in (0, 1)
# into two standard normals
def random_normals(size):

    def box_muller(u1, u2):
        R = np.sqrt(-2 * np.log(u1))
        theta = 2 * np.pi * u2

        z0 = R * np.cos(theta)
        z1 = R * np.sin(theta)

        return z0, z1

    def normal_generator(uniform_generator):
        while True:
            z0, z1 = box_muller(next(uniform_generator), next(uniform_generator))
            yield z0
            yield z1

    def random_generator(batch_size):
        while True:
            for v in np.random.uniform(size=batch_size):
                yield v

    result = np.empty(size)
    gen = normal_generator(random_generator(batch_size=min(size, 1024)))
    for i in range(size):
        result[i] = next(gen)

    return result

def get_observation():
    z = random_normals(k)
    return mu_array + sigma_sqrt @ z


return np.array([get_observation() for _ in range(nsim)])

```

```
In [4]: # Sample usage
import matplotlib.pyplot as plt
%matplotlib inline

mu = [1, 3]
sigma = [[2, 1], [1, 2]]
obs = multivariate_normal_observations(mu, sigma, nsim=5000)

plt.figure(figsize=(12, 8))
plt.hexbin(obs[:, 0], obs[:, 1], cmap=plt.cm.Reds)
plt.show()
```



**Exercise 15.6.6 (Computer Experiment).** Generate 1000 random vectors from a  $\mathcal{N}(\mu, \Sigma)$  distribution where

$$\mu = \begin{pmatrix} 3 \\ 8 \end{pmatrix}, \quad \Sigma = \begin{pmatrix} 2 & 6 \\ 6 & 2 \end{pmatrix}$$

Plot the simulation as a scatterplot. Find the distribution of  $X_2|X_1 = x_1$  using theorem 15.5. In particular, what is the formula for  $\mathbb{E}(X_2|X_1 = x_1)$ ? Plot  $\mathbb{E}(X_2|X_1 = x_1)$  on your scatterplot. Find the correlation  $\rho$  between  $X_1$  and  $X_2$ . Compare this with the sample correlations from your simulation. Find a 95% confidence interval for  $\rho$ . Estimate the covariance matrix  $\Sigma$ .

### Solution.

The provided  $\Sigma$  matrix has negative eigenvalues. We will instead use the following matrix:

$$\Sigma = \begin{pmatrix} 6 & 2 \\ 2 & 6 \end{pmatrix}$$

```
In [5]: # Generate 1000 vectors
mu = [3, 8]
sigma = [[2, 6], [6, 2]]
obs = multivariate_normal_observations(mu, sigma, nsim=1000)


# Using numpy to generate observations:
#obs = np.random.multivariate_normal(mu, sigma, size=1000)

# Using scipy to generate observations:
#obs = scipy.stats.multivariate_normal.rvs(mean=mu, cov=sigma, size=1000)

x, y = obs[:, 0], obs[:, 1]

# Plot scatterplot
plt.figure(figsize=(12, 8))
plt.scatter(x, y)
plt.show()
```





From theorem 15.5,

$$X_2|X_1 = x_1 \sim N(\mu(x_1), \Sigma(x_1))$$

where

$$\mu(x_1) = \mu_2 + \Sigma_{21}\Sigma_{11}^{-1}(x_1 - \mu_1) \quad (33)$$

$$\Sigma(x_1) = \Sigma_{22} - \Sigma_{21}\Sigma_{11}^{-1}\Sigma_{12} \quad (34)$$


Replacing the given values,

$$\mu(x_1) = 8 + 2 \cdot 6^{-1}(x_1 - 3) = \frac{1}{3}x_1 + \frac{22}{3} \quad (35)$$

$$\Sigma(x_1) = 6 - 2 \cdot 6^{-1} \cdot 2 = \frac{16}{3} \quad (36)$$

```
In [6]: # Plot scatterplot + line
        f = lambda x: x/3 + 22/3

        plt.figure(figsize=(12, 8))
        plt.scatter(x, y)
        plt.plot([-6, 10], [f(-6), f(10)], color='red')
        plt.show()
```



The correlation  $\rho$  between  $X_1$  and  $X_2$  is:

$$\rho = \frac{\text{Cov}(X_1, X_2)}{\sigma_{X_1} \sigma_{X_2}} = \frac{1}{3}$$

The estimated correlation  $\rho$  between  $X_1$  and  $X_2$  is:

$$\hat{\rho} = \frac{\sum_i (X_{1i} - \bar{X}_1)(X_{2i} - \bar{X}_2)}{s_{X_1} s_{X_2}}$$

```
In [7]: rho_hat = np.corrcoef(x, y)[0, 1]
        print("Estimated correlation: %.3f" % rho_hat)
```

Estimated correlation: 0.327

We will use the provided process to estimate a confidence interval for the correlation:

### Approximate Confidence Interval for the Correlation

1. Compute

$$\hat{\theta} = f(\hat{\rho}) = \frac{1}{2}(\log(1 + \hat{\rho}) - \log(1 - \hat{\rho}))$$

1. Compute the approximate standard error of  $\hat{\theta}$  which can be shown to be

$$\hat{\text{se}}(\hat{\theta}) = \frac{1}{\sqrt{n-3}}$$

1. An approximate  $1 - \alpha$  confidence interval for  $\theta = f(\rho)$  is

$$(a, b) \equiv \left( \hat{\theta} - \frac{z_{\alpha/2}}{\sqrt{n-3}}, \hat{\theta} + \frac{z_{\alpha/2}}{\sqrt{n-3}} \right)$$

1. Apply the inverse transformation  $f^{-1}(z)$  to get a confidence interval for  $\rho$ :

$$\left( \frac{e^{2a} - 1}{e^{2a} + 1}, \frac{e^{2b} - 1}{e^{2b} + 1} \right)$$

```
In [8]: from scipy.stats import norm

theta_hat = (np.log(1 + rho_hat) - np.log(1 - rho_hat)) / 2
se_theta_hat = 1 / np.sqrt(1000 - 3)
z = norm.ppf(0.975)
a, b = theta_hat - z * se_theta_hat, theta_hat + z * se_theta_hat
f_inv = lambda x: (np.exp(2*x) - 1) / (np.exp(2*x) + 1)
confidence_interval = (f_inv(a), f_inv(b))

print('95% confidence interval: %.3f, %.3f' % confidence_interval)
```

95% confidence interval: 0.270, 0.381

The sample covariance matrix is:

$$\hat{\Sigma} = \frac{1}{n} S = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})(X_i - \bar{X})^T$$

```
In [9]: import numpy as np
mu_hat = np.array([x.mean(), y.mean()])
xx = np.concatenate((x.reshape(-1, 1), y.reshape(-1, 1)), axis=1)
sigma_hat = (xx - mu_hat).T @ (xx - mu_hat) / 1000

sigma_hat
```

```
Out[9]: array([[6.17974344, 1.99813879],
               [1.99813879, 6.05196861]])
```

**Exercise 15.6.7 (Computer Experiment).** Generate 100 random vectors from a multivariate Normal with mean  $(0, 2)^T$  and variance

$$\begin{pmatrix} 1 & 3 \\ 3 & 1 \end{pmatrix}$$

Find a 95% confidence interval for the correlation  $\rho$ . What is the true value of  $\rho$ ?

**Solution.**

The provided matrix, yet again, has negative eigenvalues. Let's instead use:

$$\Sigma = \begin{pmatrix} 3 & 1 \\ 1 & 3 \end{pmatrix}$$

```

In [10]: # Generate 100 vectors
n = 100
mu = [0, 2]
sigma = [[3, 1], [1, 3]]
obs = multivariate_normal_observations(mu, sigma, nsim=n)
x, y = obs[:, 0], obs[:, 1]

In [11]: # Find 95% confidence interval
from scipy.stats import norm

rho_hat = np.corrcoef(x, y)[0, 1]
theta_hat = (np.log(1 + rho_hat) - np.log(1 - rho_hat)) / 2
se_theta_hat = 1 / np.sqrt(n - 3)
z = norm.ppf(0.975)
a, b = theta_hat - z * se_theta_hat, theta_hat + z * se_theta_hat
f_inv = lambda x: (np.exp(2*x) - 1) / (np.exp(2*x) + 1)
confidence_interval = (f_inv(a), f_inv(b))

print('95% confidence interval: %.3f, %.3f' % confidence_interval)

```

95% confidence interval: 0.054, 0.424

True value of  $\rho$ :

$$\rho = \frac{\text{Cov}(X_1, X_2)}{\sigma_{X_1}\sigma_{X_2}} = \frac{1}{3}$$